

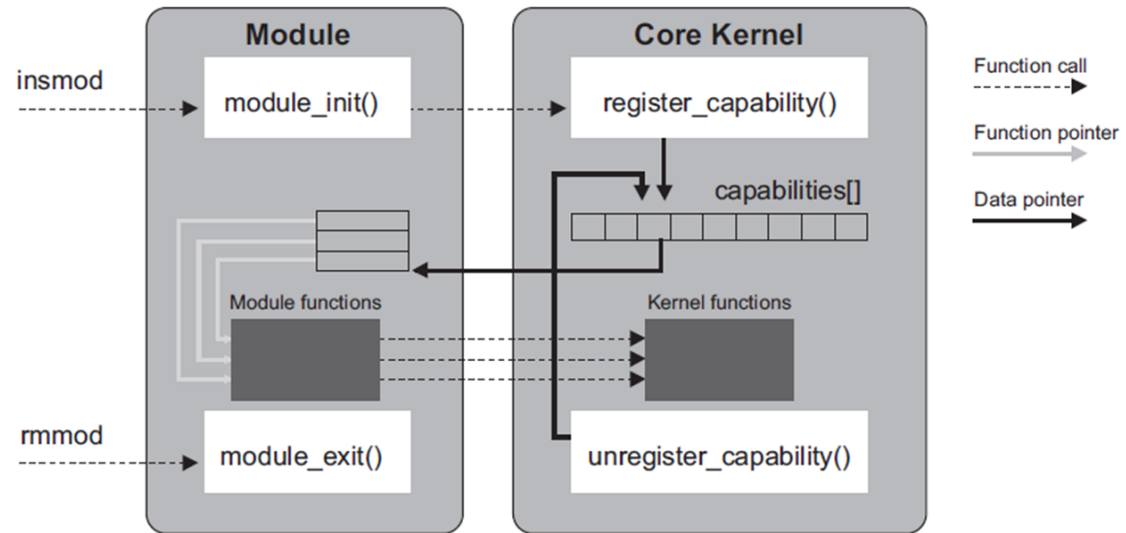
# Operating Systems

## Recitation 11

# Plan

- Controlling driver (ioctl)
- Intercepting system call
- Debug-FS
- Keyboard sniffer

Last week we saw a simple character device driver. Let's add “encryption upon demand” capability.



So far —

Driver registered operations, kernel invoked those operations when needed.

Now —

We want to alter driver behavior, i.e. to “talk” to driver not to the device it controls.

# file\_operations ioctl()

```
int (*ioctl)(struct file *filp,  
             unsigned int cmd,  
             unsigned long arg);
```

- For special commands not covered by existing functions (read, write,...)
  - Considered bad programming to use it otherwise
- ~~inode – device inode file~~ (Deprecated)
- filp – open device file
- cmd – command number
- arg – optional general purpose
- Returns – depends on implementation, <0 on error

## The “TO DO” list:

1. Define the operation - flip “with encryption” flag.

We need to create a new header file

2. Write code that gets that command and reacts on it.

Implement and register certain `*ioctl*` function in driver.

3. Write user space code that raises/drops the flag.

Use the header file from #1

# CHARDEV-2

Device driver changes:

- create header file, predefine MAJOR\_NUM
- add device\_ioctl function
- register it in file\_operations structure

# CHARDEV-2

In user space:

- create `user_chardev.c`

The only thing it does:

invokes `ioctl` function for our device file,  
sets encryption flag to 1.

# The flow

```
[ 9228.487115] device_open(d5d85780)
[ 9227.487150] device_write(d5d85780,8)
[ 9227.487161] device_release(f308cb58,d5d85780)
[ 9230.357332] device_open(f4c2fd80)
[ 9230.357368] device_read(f4c2fd80,131072) - operation not supported yet (last written - Hello OS)
[ 9230.358490] device_release(f308cb58,f4c2fd80)
[ 9320.436419] device_open(d7639840)
[ 9320.436438] chardev, ioctl: setting encryption flag to 1
[ 9320.436444] device_release(f308cb58,d7639840)
[ 9335.459575] device_open(d7639840)
[ 9335.459616] device_write(d7639840,8)
[ 9335.459627] device_release(f308cb58,d7639840)
[ 9338.669655] device_open(f53dfc00)
[ 9338.669693] device_read(f53dfc00,131072) - operation not supported yet (last written - Ifmmp!PT)
[ 9338.670776] device_release(f308cb58,f53dfc00)
```

Example – CHARDE2



# Intercepting system call

- Kernel exports system calls table  
`cat /proc/kallsyms | grep sys_call_table`
- It is an array of pointers to functions
- Let's substitute one of them

Problems:

It is hidden.

It is write protected.

# Intercepting system call

Solution:

Take an address of a known function (sys\_close)

Iterate through kernel memory,

Check

- if current position can be considered

- as the sys\_call\_table,

- then the certain entry contains

- the pointer to that known function

# Intercepting system call

Once `sys_call_table` found,

Access control register `cr0` (in CPU)

Make the page writable

Preserve the pointer to the original `syscall` function

Assign the pointer to our new function

Restore read-only mode

# Intercepting system call

Upon system call:

Do additional things (printf)

Invoke original system call

# Intercepting system call

Cleaning up

Access control register cr0

Make page writable

Restore the pointer to the original syscall function

Restore read-only mode

Example - Interceptor

# Debug-FS

debugfs is a simple memory-based filesystem, designed specifically to debug Linux kernel code

Setting up:

```
mount -t debugfs nodev /sys/kernel/debug
```

(Probably done already)

# Debug-FS

## The typical sequence

On init

`debugfs_create_dir`

`debugfs_create_file`

On some event:

`pr_debug`

On cleanup

`debugfs_remove_recursive`

# Keyboard sniffer

## Problem:

- Standard file\_operations are not enough
- Too much information to put into a standard log with printk

## Solution:

Call register\_keyboard\_notifier to register certain callbacks.

Use debugfs to store all the keys pressed



# Keyboard sniffer

Linux uses **notification chains**.

These notifier chains work in a **Publish-Subscribe model**.  
This model is more effective when compared to polling or the request-reply model.

Descriptor structure -  
`struct notifier_block`

Example - Keysniffer

# Keyboard sniffer

register\_keyboard\_notifier  
start sniffing

unregister\_keyboard\_notifier  
remove the sniffer

Example - keysniffer